

Participants: Mariuxi Leon, Cameron Milligan, Jorge Gonzalez

Data Analysis

The data consist of the number of people present at different times in gym and also includes additional factor. Using this dataset interesting insights into attendance of the gym shall be derived. Also a predictive model shall be developed which will, predict the number of people attending the gym given the values of other factors.

Basic to do list

Specify the type of data analytic question (e.g. exploration, association causality) before touching the data

1. How attendance is affected by time of the week
2. How attendance is affected by time of the temperature
3. How attendance is affected by the start of the semester
4. Create a model predicting the number of people attending the gym

Define the metric for success before beginning

1. Logical reasoning shall match
2. Logical reasoning shall match
3. Logical reasoning shall match
4. 90% accuracy

Business Application

This data from university gym. This data consists of information about attendance of the gym and other factors taken at the same timestamp. Using this data a predictive model can be developed which will predict the attendance of the gym which might be useful for the gym owner.

```
In [1]: import datetime
import numpy as np
import pandas as pd
import seaborn as sb

# We will use matplotlib to plot figures
import matplotlib.pyplot as plt
%matplotlib inline

# For regression analysis we will use the statsmodels package
import statsmodels.api as sm
import statsmodels.formula.api as smf

# For visual inspection of the regression models
from statsmodels.graphics.regressionplots import plot_regress_exog, plot_fit, plot_leverage_resid2, influence_plot

# This function will help us to create ordinal variables
from pandas.api.types import CategoricalDtype

/Users/cammilligan/anaconda/lib/python3.6/site-packages/statsmodels/compat/pandas.py:56: FutureWarning: The pandas.core.datetools module is deprecated and will be removed in a future version. Please use the pandas.tseries module instead.
    from pandas.core import datetools
```

```
In [2]: df = pd.read_csv('https://gist.githubusercontent.com/cameronmilligan/d4b77b01e5fd03fd1d1b26f10f9d0c9c/raw/c807fd88a46698170151e283c749de44fde5e7e5/data.csv',)
df = df[["number_people", "timestamp", "day_of_week", "is_weekend", "is_holiday", "temperature", "is_start_of_semester", "is_during_semester", "date", "month", "hour"]]
#converting to celcius
df.temperature = (df.temperature-32)/1.8
df_copy = df.copy()
```

```
In [3]: df.temperature.describe()
```

```
Out[3]: count      62184.000000
mean          14.753949
std           3.509109
min           3.411111
25%          12.777778
50%          14.633333
75%          16.822222
max           30.650000
Name: temperature, dtype: float64
```

In [4]: `df.head(1)`

Out[4]:

	number_people	timestamp	day_of_week	is_weekend	is_holiday	temperature	is_still
0	37	61211	4	0	0	22.088889	0

In [5]: `df.shape`

Out[5]: (62184, 11)

In [6]: `df.describe()`

Out[6]:

	number_people	timestamp	day_of_week	is_weekend	is_holiday	temperature
count	62184.000000	62184.000000	62184.000000	62184.000000	62184.000000	62184.000000
mean	29.072543	45799.437958	2.982504	0.282870	0.002573	14.7539
std	22.689026	24211.275891	1.996825	0.450398	0.050660	3.50910
min	0.000000	0.000000	0.000000	0.000000	0.000000	3.4111
25%	9.000000	26624.000000	1.000000	0.000000	0.000000	12.777
50%	28.000000	46522.500000	3.000000	0.000000	0.000000	14.633
75%	43.000000	66612.000000	5.000000	1.000000	0.000000	16.822
max	145.000000	86399.000000	6.000000	1.000000	1.000000	30.650

The dataset seems to be clean because:

Number of rows for all the data columns are same

Maximum number of people = 145 and min = 0 , looks reasonable

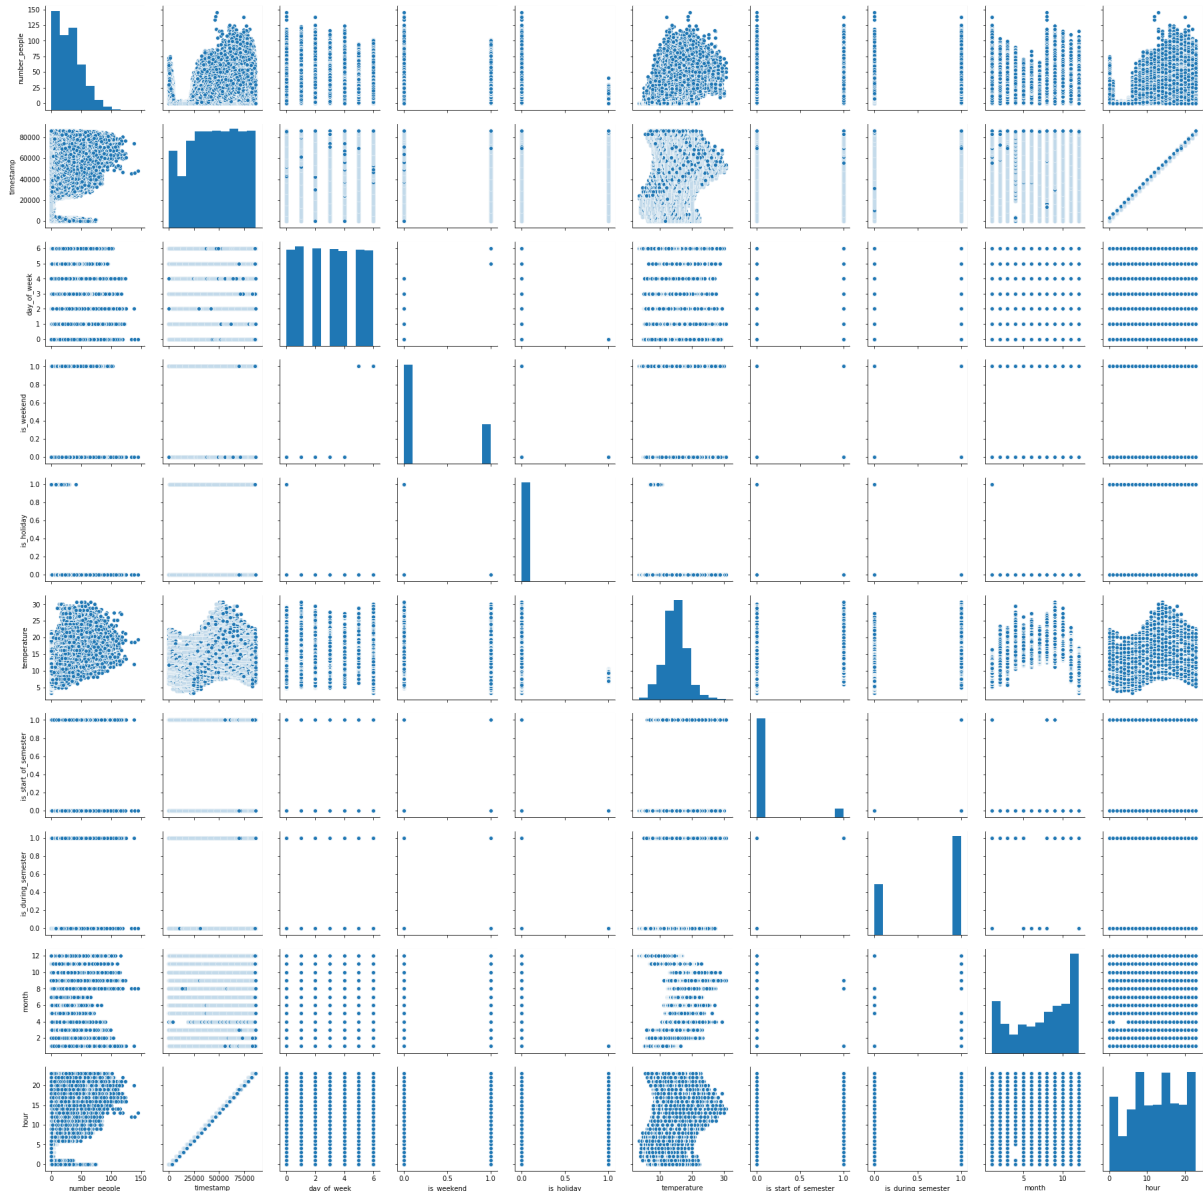
Maximum timestampvalue = 86399 which is close to 24 hrs which is reasonable

Max day of the week = 6 and min = 0 , seems reasonable

Below we pair plot, to get pictorial overview of the complete data. Pairplot shows histograms of the columns in diagonal of the matrix and pairwise scatter plot of the data.

```
In [7]: sb.pairplot(df)
```


```
Out[7]: <seaborn.axisgrid.PairGrid at 0x119874d30>
```



From above pairplot few things can be implied

- 1. Number of people attending gym and temperatures are showing gaussian distribution which is good thing**
- 2. There is some kind of relationship existing between temperatures and number of people attending gym**

```
In [8]: corrmat = df.corr()
f, ax = plt.subplots(figsize=(9, 9))
# Draw the heatmap using seaborn
sb.heatmap(corrmat, square=False, annot=True)
plt.show()
```



Some of the correlations above can be ignored such as day_of_week and is_weekend being correlated.

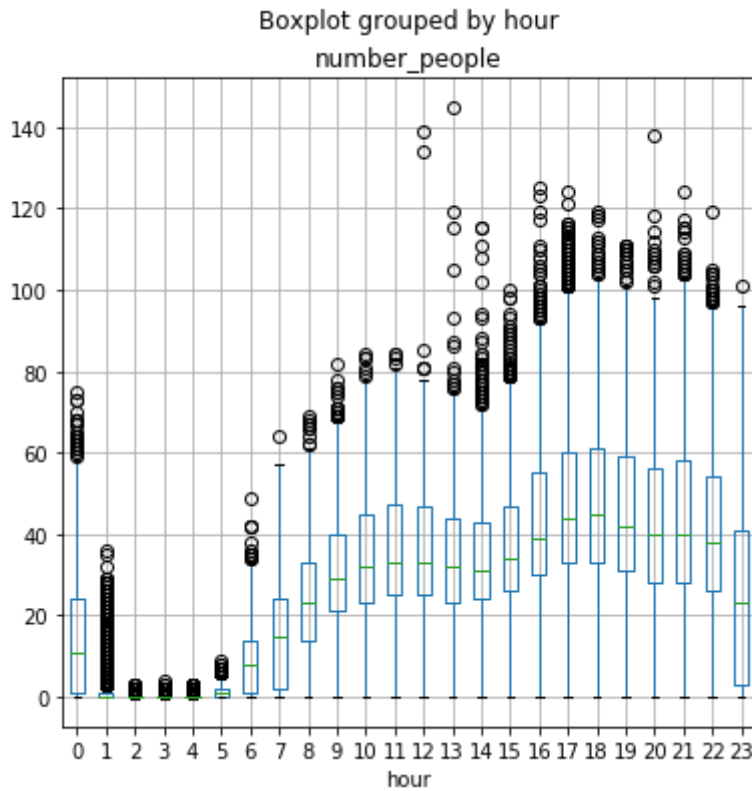
Some of the correlations that stand out are the hour, day_of_week/is_weekend, temperature and whether or not it is the start of the semester

Analysis: Hour vs Number of Gym Attendees

Hypothesis: A greater number of people are attending the gym later in the day.

```
In [9]: df.boxplot(column="number_people", by= "hour", figsize= (6,6))
```


```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x127c92cc0>
```



```
In [10]: sb.barplot(x='hour',y='number_people',data=df)
```

```
/Users/cammilligan/anaconda/lib/python3.6/site-packages/scipy/stats/stats.py:1626: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.
    return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1294ce518>
```



Statement: "higher number of people are attending the gym at later hours" appears true from above graph

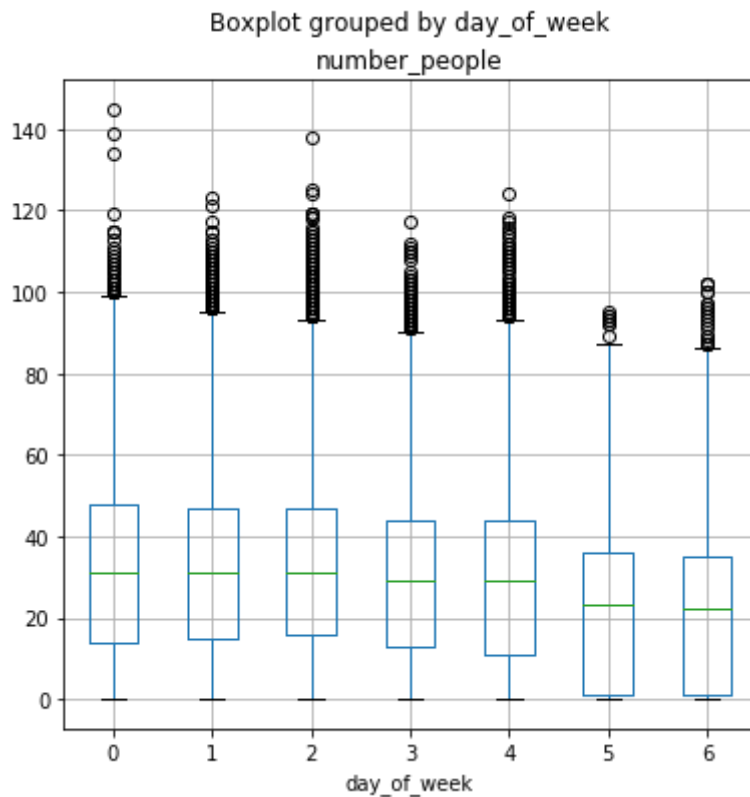
Analysis: Day of the Week vs Number of Gym Attendees

From pairplot there is negative correlation between day_of_week and number of people going to gym.

Hypothesis: A greater number of people attending the gym on day = 0 i.e monday compared day = 6 i.e sunday.


```
In [11]: df.boxplot(column="number_people", by= "day_of_week", figsize= (6,6))
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x129704d30>
```



```
In [12]: sb.barplot(x='day_of_week',y='number_people',data=df)
```


```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x1299f61d0>
```



Statement: Gym attendance steadily declines from Monday through Sunday.


```
In [13]: df.boxplot(column="number_people", by= "is_weekend", figsize= (6,6))
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x129c26be0>
```



The above statement can be analyzed in another angle by bucketing the data further and comparing weekdays versus weekend days. It is clear more people attend the gym during the week than on weekends.

Analysis: Temperature vs Number of Gym Attendees

From correlation plot, there exists the positive correlation between number of people going to gym and temperature.


Hypothesis: The higher the temperature the lower the gym attendance.

Temperature data is continuous in nature, so using a histogram is a better way to visualize it compared to boxplots or bar charts

```
In [14]: Bins = []
for i in range(0,35,5):
    people_count = 0
    for index, row in df.iterrows():
        if (row['temperature'] >= i and row['temperature'] < i+5):
            people_count = people_count + row['number_people']
    Bins.append((people_count))


sb.barplot(list(range(0,35,5)),Bins)
```

Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x129fa9198>




```
In [15]: df.plot(kind='bar',x='temperature',y='number_people', figsize=(12,9))
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x12a7de470>
```




```
In [16]: df.hist(column='temperature',bins=40)
```

```
Out[16]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x12a67e3c8
>]],
      dtype=object)
```



```
In [17]: sb.distplot(df['temperature'], kde=False, rug=True)
```


```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x12bae64e0>
```



Even though from the distribution it seems like more people attend gym in warm weather its not true as most number of days weather remains in range 50 to 70. So the analysis that more people attend gym during 50 to 70 temperature range is not entirely correct.


```
In [18]: df.boxplot(column="number_people", by= "is_holiday", figsize= (6,6))
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x14469b6a0>
```




```
In [19]: df.boxplot(column="number_people", by= "is_start_of_semester", figsize=(6,6))
```

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x14541d240>
```




```
In [20]: df.boxplot(column="number_people", by= "is_during_semester", figsize= (6,6))
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x14559f7b8>
```



```
In [21]: df.boxplot(column="number_people", by= "month", figsize= (8,8))
#could this basically be the same information as is_during_semester and
is_start_of_semester? concerns about overfitting maybe? Worth at least
discussing in the assignment
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x12bc1f710>
```



```
In [22]: df.dtypes
```


```
Out[22]: number_people      int64
timestamp      int64
day_of_week     int64
is_weekend      int64
is_holiday      int64
temperature    float64
is_start_of_semester  int64
is_during_semester  int64
date           object
month          int64
hour           int64
dtype: object
```



```
In [23]: #A lot of the features are categorical, so they have to be changed to the category datatype
df["day_of_week"] = df["day_of_week"].astype('category')
df["is_weekend"] = df["is_weekend"].astype('category')
df["is_holiday"] = df["is_holiday"].astype('category')
df["is_start_of_semester"] = df["is_start_of_semester"].astype('category')
df["is_during_semester"] = df["is_during_semester"].astype('category')
df["month"] = df["month"].astype('category')
df["hour"] = df["hour"].astype('category')
```

```
In [24]: # Initialize and fit the model
m1 = smf.ols(formula='number_people ~ day_of_week + is_weekend + is_holiday + temperature + is_start_of_semester + is_during_semester + month + hour', data=df)
m1 = m1.fit()
df['number_people_predicted'] = m1.predict(df)
df['number_people_residuals'] = df['number_people_predicted'] - df['number_people']
# Residuals mean
print("residuals mean: ",df.number_people_residuals.mean())
print("residuals mean: ",m1.rsquared)
# Plot histogram of the residuals
plt.hist(df.number_people_residuals,bins=40)
plt.xlabel('number of people residuals')
plt.title('Including is_weekend')
plt.show()
```

```
residuals mean: -1.002208736687098e-13
residuals mean: 0.6265773532323791
```



```
In [25]: print(ml.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable:          number_people    R-squared:
0.627
Model:                  OLS              Adj. R-squared:
0.626
Method:                 Least Squares    F-statistic:
2370.
Date:                   Mon, 16 Sep 2019  Prob (F-statistic):
0.00
Time:                   18:25:32          Log-Likelihood:           -2.5
174e+05
No. Observations:      62184             AIC:                       5.
036e+05
Df Residuals:          62139             BIC:                       5.
040e+05
Df Model:               44
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t
[0.025 0.975]				

Intercept	3.7228	0.451	8.258	0.000
2.839 4.606				
day_of_week[T.1]	-0.5321	0.208	-2.556	0.011
-0.940 -0.124				
day_of_week[T.2]	-1.0535	0.209	-5.036	0.000
-1.464 -0.643				
day_of_week[T.3]	-2.8412	0.210	-13.561	0.000
-3.252 -2.431				
day_of_week[T.4]	-3.5789	0.211	-16.993	0.000
-3.992 -3.166				
day_of_week[T.5]	-2.9220	0.121	-24.162	0.000
-3.159 -2.685				
day_of_week[T.6]	-3.7137	0.121	-30.666	0.000
-3.951 -3.476				
is_weekend[T.1]	-6.6357	0.121	-54.649	0.000
-6.874 -6.398				
is_holiday[T.1]	-17.3138	1.135	-15.249	0.000
-19.539 -15.088				
is_start_of_semester[T.1]	3.1802	0.279	11.385	0.000
2.633 3.728				
is_during_semester[T.1]	14.8879	0.230	64.765	0.000
14.437 15.339				
month[T.2]	-4.6966	0.388	-12.111	0.000
-5.457 -3.937				
month[T.3]	-11.6965	0.377	-30.989	0.000
-12.436 -10.957				
month[T.4]	-8.0203	0.416	-19.287	0.000
-8.835 -7.205				
month[T.5]	-12.5490	0.342	-36.659	0.000
-13.220 -11.878				
month[T.6]	-7.2086	0.357	-20.181	0.000
-7.909 -6.508				

month[T.7]		-5.5862	0.346	-16.166	0.000
-6.264	-4.909				
month[T.8]		-8.0697	0.319	-25.333	0.000
-8.694	-7.445				
month[T.9]		-8.5284	0.352	-24.195	0.000
-9.219	-7.838				
month[T.10]		-12.0302	0.377	-31.928	0.000
-12.769	-11.292				
month[T.11]		-13.8252	0.355	-38.935	0.000
-14.521	-13.129				
month[T.12]		-12.7204	0.306	-41.602	0.000
-13.320	-12.121				
hour[T.1]		-10.3511	0.416	-24.888	0.000
-11.166	-9.536				
hour[T.2]		-11.5346	0.438	-26.305	0.000
-12.394	-10.675				
hour[T.3]		-11.3915	0.439	-25.950	0.000
-12.252	-10.531				
hour[T.4]		-10.5569	0.447	-23.610	0.000
-11.433	-9.681				
hour[T.5]		-12.3209	0.379	-32.489	0.000
-13.064	-11.578				
hour[T.6]		-5.1348	0.367	-13.988	0.000
-5.854	-4.415				
hour[T.7]		1.1104	0.366	3.037	0.002
0.394	1.827				
hour[T.8]		9.2578	0.365	25.350	0.000
8.542	9.974				
hour[T.9]		14.7418	0.365	40.335	0.000
14.025	15.458				
hour[T.10]		17.5772	0.367	47.946	0.000
16.859	18.296				
hour[T.11]		19.4239	0.369	52.689	0.000
18.701	20.146				
hour[T.12]		18.6265	0.370	50.276	0.000
17.900	19.353				
hour[T.13]		16.5558	0.373	44.413	0.000
15.825	17.286				
hour[T.14]		15.7203	0.375	41.962	0.000
14.986	16.455				
hour[T.15]		19.0739	0.374	50.964	0.000
18.340	19.808				
hour[T.16]		25.1287	0.372	67.577	0.000
24.400	25.858				
hour[T.17]		30.4123	0.367	82.938	0.000
29.694	31.131				
hour[T.18]		31.5467	0.367	85.999	0.000
30.828	32.266				
hour[T.19]		28.5130	0.366	77.897	0.000
27.796	29.230				
hour[T.20]		25.8457	0.366	70.589	0.000
25.128	26.563				
hour[T.21]		27.4416	0.366	74.920	0.000
26.724	28.159				
hour[T.22]		25.3289	0.366	69.177	0.000
24.611	26.047				
hour[T.23]		11.2920	0.363	31.117	0.000

```

10.581      12.003
temperature      1.0314      0.025      41.704      0.000
0.983      1.080
=====
=====
Omnibus:      3845.022      Durbin-Watson:
0.234
Prob(Omnibus):      0.000      Jarque-Bera (JB):      7
366.758
Skew:      0.451      Prob(JB):
0.00
Kurtosis:      4.425      Cond. No.
2.25e+15
=====
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.
[2] The smallest eigenvalue is 2.84e-24. This might indicate that there
are
strong multicollinearity problems or that the design matrix is singula
r.

```

In [26]: `df.head(1)`

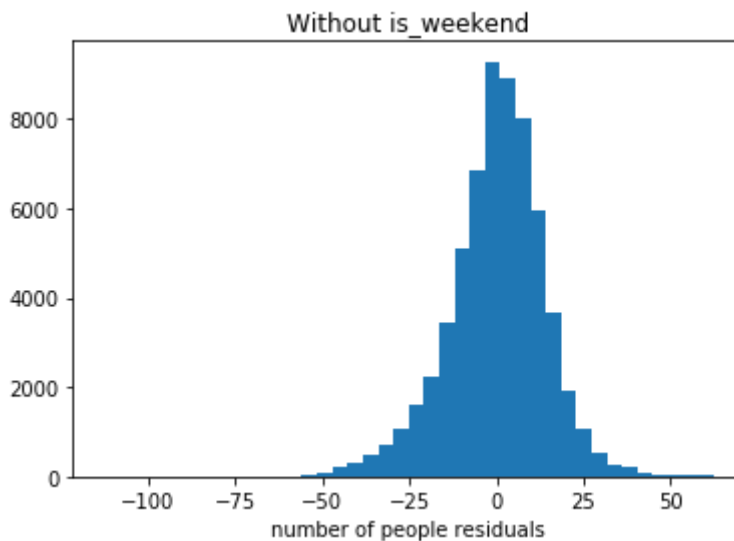
Out[26]:

	number_people	timestamp	day_of_week	is_weekend	is_holiday	temperature	is_sta
0	37	61211	4	0	0	22.088889	0

```
In [27]: # Initialize and fit the model
params = ['day_of_week' , 'is_holiday' , 'temperature' , 'is_start_of_se
mester' , 'is_during_semester' , 'month' , 'hour']
m2 = smf.ols(formula='number_people ~ day_of_week + is_holiday + tempera
ture + is_start_of_semester + is_during_semester + month + hour', data=d
f)
m2 = m2.fit()

df['number_people_predicted'] = m2.predict(df)
df['number_people_redisuals'] = df['number_people_predicted'] - df['numb
er_people']
# Residuals mean
print("residuals mean: ",df.number_people_redisuals.mean())
print("residuals mean: ",m2.rsquared)
# Plot histogram of the residuals
plt.hist(df.number_people_redisuals,bins=40)
plt.xlabel('number of people residuals')
plt.title('Without is_weekend')
plt.show()
```

```
residuals mean: -2.516511235711655e-13
residuals mean: 0.626577353232379
```



```
In [28]: print(m2.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable:          number_people    R-squared:
0.627
Model:                  OLS             Adj. R-squared:
0.626
Method:                Least Squares    F-statistic:
2370.
Date:                  Mon, 16 Sep 2019  Prob (F-statistic):
0.00
Time:                  18:25:34         Log-Likelihood:          -2.5
174e+05
No. Observations:      62184           AIC:                      5.
036e+05
Df Residuals:          62139           BIC:                      5.
040e+05
Df Model:              44
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t
[0.025 0.975]				

Intercept	3.7228	0.451	8.258	0.000
2.839 4.606				
day_of_week[T.1]	-0.5321	0.208	-2.556	0.011
-0.940 -0.124				
day_of_week[T.2]	-1.0535	0.209	-5.036	0.000
-1.464 -0.643				
day_of_week[T.3]	-2.8412	0.210	-13.561	0.000
-3.252 -2.431				
day_of_week[T.4]	-3.5789	0.211	-16.993	0.000
-3.992 -3.166				
day_of_week[T.5]	-9.5577	0.210	-45.528	0.000
-9.969 -9.146				
day_of_week[T.6]	-10.3494	0.210	-49.230	0.000
-10.761 -9.937				
is_holiday[T.1]	-17.3138	1.135	-15.249	0.000
-19.539 -15.088				
is_start_of_semester[T.1]	3.1802	0.279	11.385	0.000
2.633 3.728				
is_during_semester[T.1]	14.8879	0.230	64.765	0.000
14.437 15.339				
month[T.2]	-4.6966	0.388	-12.111	0.000
-5.457 -3.937				
month[T.3]	-11.6965	0.377	-30.989	0.000
-12.436 -10.957				
month[T.4]	-8.0203	0.416	-19.287	0.000
-8.835 -7.205				
month[T.5]	-12.5490	0.342	-36.659	0.000
-13.220 -11.878				
month[T.6]	-7.2086	0.357	-20.181	0.000
-7.909 -6.508				
month[T.7]	-5.5862	0.346	-16.166	0.000
-6.264 -4.909				

month[T.8]		-8.0697	0.319	-25.333	0.000
-8.694	-7.445				
month[T.9]		-8.5284	0.352	-24.195	0.000
-9.219	-7.838				
month[T.10]		-12.0302	0.377	-31.928	0.000
-12.769	-11.292				
month[T.11]		-13.8252	0.355	-38.935	0.000
-14.521	-13.129				
month[T.12]		-12.7204	0.306	-41.602	0.000
-13.320	-12.121				
hour[T.1]		-10.3511	0.416	-24.888	0.000
-11.166	-9.536				
hour[T.2]		-11.5346	0.438	-26.305	0.000
-12.394	-10.675				
hour[T.3]		-11.3915	0.439	-25.950	0.000
-12.252	-10.531				
hour[T.4]		-10.5569	0.447	-23.610	0.000
-11.433	-9.681				
hour[T.5]		-12.3209	0.379	-32.489	0.000
-13.064	-11.578				
hour[T.6]		-5.1348	0.367	-13.988	0.000
-5.854	-4.415				
hour[T.7]		1.1104	0.366	3.037	0.002
0.394	1.827				
hour[T.8]		9.2578	0.365	25.350	0.000
8.542	9.974				
hour[T.9]		14.7418	0.365	40.335	0.000
14.025	15.458				
hour[T.10]		17.5772	0.367	47.946	0.000
16.859	18.296				
hour[T.11]		19.4239	0.369	52.689	0.000
18.701	20.146				
hour[T.12]		18.6265	0.370	50.276	0.000
17.900	19.353				
hour[T.13]		16.5558	0.373	44.413	0.000
15.825	17.286				
hour[T.14]		15.7203	0.375	41.962	0.000
14.986	16.455				
hour[T.15]		19.0739	0.374	50.964	0.000
18.340	19.808				
hour[T.16]		25.1287	0.372	67.577	0.000
24.400	25.858				
hour[T.17]		30.4123	0.367	82.938	0.000
29.694	31.131				
hour[T.18]		31.5467	0.367	85.999	0.000
30.828	32.266				
hour[T.19]		28.5130	0.366	77.897	0.000
27.796	29.230				
hour[T.20]		25.8457	0.366	70.589	0.000
25.128	26.563				
hour[T.21]		27.4416	0.366	74.920	0.000
26.724	28.159				
hour[T.22]		25.3289	0.366	69.177	0.000
24.611	26.047				
hour[T.23]		11.2920	0.363	31.117	0.000
10.581	12.003				
temperature		1.0314	0.025	41.704	0.000

```

0.983      1.080
=====
=====
Omnibus:      3845.022    Durbin-Watson:
0.234
Prob(Omnibus):      0.000    Jarque-Bera (JB):      7
366.758
Skew:      0.451    Prob(JB):
0.00
Kurtosis:      4.425    Cond. No.
350.
=====
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Analysis

At start of semester people tend to go to gym more

```

In [29]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import SGDClassifier,SGDRegressor

```

```

/Users/cammilligan/anaconda/lib/python3.6/site-packages/sklearn/ensembl
e/weight_boosting.py:29: DeprecationWarning: numpy.core.umath_tests is
an internal NumPy module and should not be imported. It will be removed
in a future NumPy release.

```

```

from numpy.core.umath_tests import inner1d

```

```
In [30]: df2 = df_copy.copy()
y = df2['number_people'].values
df2_X = df2.drop(['number_people', 'date', 'timestamp', 'is_weekend'], axis=
1)
X = df2_X.values
df2.head(2)
```

Out[30]:

	number_people	timestamp	day_of_week	is_weekend	is_holiday	temperature	is_sta
0	37	61211	4	0	0	22.088889	0
1	45	62414	4	0	0	22.088889	0

```
In [31]: Xtrain, Xtest, ytrain, ytest = train_test_split(X, y, test_size=0.40)
```

```
In [32]: # Scale the data to be between -1 and 1
scaler = StandardScaler()
scaler.fit(Xtrain)
Xtrain = scaler.transform(Xtrain)
Xtest = scaler.transform(Xtest)
```

```
In [33]: sdg = SGDRegressor()
sdg.fit(Xtrain, ytrain)
y_val_1 = sdg.predict(Xtest)
print(sdg.score(Xtest, ytest))
```

0.505364392240245

```
In [34]: radm = RandomForestRegressor(n_estimators=100)
radm.fit(Xtrain, ytrain)

print(radm.score(Xtest, ytest))
```

0.912177349613403

Analysis

The predictive model can predict number of people attending the gym with 91% accuracy

```
In [35]: indices = np.argsort(radm.feature_importances_)[::-1]

# Print the feature ranking
print('Feature ranking:')

for f in range(df2_X.shape[1]):
    print('%d. feature %d %s (%f)' % (f+1, indices[f], df2_X.columns[indices[f]],
                                     radm.feature_importances_[indices[f]]))
```


Feature ranking:

```
1. feature 6 hour (0.519484)
2. feature 2 temperature (0.179122)
3. feature 4 is_during_semester (0.110511)
4. feature 0 day_of_week (0.093581)
5. feature 5 month (0.084394)
6. feature 3 is_start_of_semester (0.012813)
7. feature 1 is_holiday (0.000095)
```

```
In [36]: df2['number_people_predicted'] = radm.predict(X)
```

```
In [37]: df2['number_people_redisuals'] = df2['number_people_predicted'] - df2['number_people']
# Residuals mean
print("residuals mean: ",df2.number_people_redisuals.mean())
# Plot histogram of the residuals
plt.hist(df2.number_people_redisuals,bins=40)
plt.xlabel('number of people residuals')
plt.title('Random Forest Residuals')
plt.show()
```

residuals mean: 15.401467539222255



Analysis

Below is the list of the factors affecting the number of people attending gym ranked from top to bottom in decreasing importance manner:

1. Hour
2. Temperature
3. Is during semester
4. Day of week
5. Month
6. Is start of semester
7. Is holiday